

## I Am Not Doing EDA In This Notebook File, Because EDA Is Already Given And Also I Have Done EDA On Another Notebook File, This Is A Notebook For Questions Solution

In [1]:

```
# import important libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings as warning
warning.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
```

In [2]:

```
# read dataset from external storage
dff = pd.read_csv('C:/Users/mridh/Downloads/EDA+ML-Final Project/nyc_taxi_trip_duration.csv')
```

## First I Am Taking 50000 Sample Data Randomly From Given Dataset For Model Building And Prediction, Because Given Dataset Is To Large For Model Building & Prediction

In [3]:

```
# creating sample dataset
# data = dff.sample(n = 50000)
```

In [4]:

```
# making csv file from sample dataset
# data.to_csv('nyc_taxi_trip_sample_data.csv', index = False)
```

In [5]:

```
# import new sample dataset
df = pd.read_csv('nyc_taxi_trip_sample_data.csv')
df.head()
```

Out[5]:

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude
0	id0045738	1	2016-01-12 15:25:16	2016-01-12 15:36:29	1	-73.967026
1	id2483471	2	2016-03-31 11:43:59	2016-03-31 12:01:09	1	-73.971046
2	id0697786	1	2016-01-11 19:24:16	2016-01-11 19:31:12	2	-73.998093
3	id0550679	2	2016-02-09 18:07:40	2016-02-09 18:11:28	1	-73.990631
4	id0786941	2	2016-02-26 17:46:13	2016-02-26 17:58:54	1	-73.959160

In [6]:

```
# check shape of dataframe
df.shape
```

Out[6]:

(50000, 11)

## Handle Missing And Duplicate Values

In [7]:

```
# check null values
df.isnull().sum()
```

Out[7]:

```
id                0
vendor_id         0
pickup_datetime   0
dropoff_datetime  0
passenger_count   0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
store_and_fwd_flag 0
trip_duration     0
dtype: int64
```

In [8]:

```
# check duplicate values
df.duplicated().sum()
```

Out[8]:

0

**We can see there are no null and duplicate values**

In [9]:

```
# converting strings to datetime features
df['pickup_datetime'] = pd.to_datetime(df.pickup_datetime)
df['dropoff_datetime'] = pd.to_datetime(df.dropoff_datetime)

df['day_of_week'] = df['pickup_datetime'].dt.weekday
df['hour_of_day'] = df['pickup_datetime'].dt.hour
df['year'] = df['pickup_datetime'].dt.year
df['month'] = df['pickup_datetime'].dt.month
df['day_of_month'] = df['pickup_datetime'].dt.day
```

In [10]:

```
df = df.drop(['id', 'pickup_datetime', 'dropoff_datetime'], axis = 1)
```

In [11]:

```
# check information of the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   vendor_id             50000 non-null  int64
1   passenger_count       50000 non-null  int64
2   pickup_longitude      50000 non-null  float64
3   pickup_latitude       50000 non-null  float64
4   dropoff_longitude     50000 non-null  float64
5   dropoff_latitude      50000 non-null  float64
6   store_and_fwd_flag    50000 non-null  object
7   trip_duration         50000 non-null  int64
8   day_of_week           50000 non-null  int64
9   hour_of_day           50000 non-null  int64
10  year                  50000 non-null  int64
11  month                 50000 non-null  int64
12  day_of_month          50000 non-null  int64
dtypes: float64(4), int64(8), object(1)
memory usage: 5.0+ MB
```

## Handle Incorrect Values

In [12]:

```
# check value counts in vendor_id column
df.vendor_id.value_counts()
```

Out[12]:

```
2    26821
1    23179
Name: vendor_id, dtype: int64
```

In [13]:

```
# check value counts in passenger_count column
df.passenger_count.value_counts()
```

Out[13]:

```
1    35539
2     7164
5     2687
3     2004
6     1617
4      983
0         6
Name: passenger_count, dtype: int64
```

**We can see there is only 6 data for passenger count 0, so we have to drop it because it can mislead our model**

In [14]:

```
# drop outliers from passenger count column
df = df[(df['passenger_count']>=1) & (df['passenger_count']<=6)]
```

In [15]:

```
# check value counts in pickup_longitude column
df.pickup_longitude.value_counts()
```

Out[15]:

```
-73.982201    27
-73.991508    26
-73.982178    25
-73.982338    25
-73.982208    24
..
-73.960274     1
-73.965424     1
-73.945923     1
-73.864182     1
-73.918388     1
Name: pickup_longitude, Length: 10915, dtype: int64
```

In [16]:

```
# check value counts in pickup_latitude column
df.pickup_latitude.value_counts()
```

Out[16]:

```
40.774052    18
40.773998    16
40.774101    15
40.754711    15
40.750259    14
..
40.751415     1
40.750885     1
40.760765     1
40.737976     1
40.767971     1
Name: pickup_latitude, Length: 20423, dtype: int64
```

In [17]:

```
# check value counts in dropoff_longitude column
df.dropoff_longitude.value_counts()
```

Out[17]:

```
-73.991379    23
-73.981972    22
-73.981987    22
-73.982246    21
-73.982140    21
..
-73.906815     1
-74.041252     1
-74.005432     1
-73.939430     1
-74.002556     1
Name: dropoff_longitude, Length: 12716, dtype: int64
```

In [18]:

```
# check value counts in dropoff_latitude column
df.dropoff_latitude.value_counts()
```

Out[18]:

```
40.774311    14
40.774281    14
40.764610    13
40.762032    13
40.764130    13
..
40.706917     1
40.697231     1
40.735798     1
40.778355     1
40.727200     1
Name: dropoff_latitude, Length: 22763, dtype: int64
```

In [19]:

```
# check value counts in store_and_fwd_flag column
df.store_and_fwd_flag.value_counts()
```

Out[19]:

```
N    49705
Y      289
Name: store_and_fwd_flag, dtype: int64
```

In [20]:

```
# check value counts in trip_duration column
df.trip_duration.value_counts()
```

Out[20]:

```
422    77
321    68
461    68
416    67
379    67
..
6090    1
3090    1
3445    1
4307    1
2948    1
Name: trip_duration, Length: 3516, dtype: int64
```

In [21]:

```
# check value counts in day_of_week column
df.day_of_week.value_counts()
```

Out[21]:

```
4    7682
5    7469
3    7427
2    7240
1    6942
6    6851
0    6383
Name: day_of_week, dtype: int64
```

In [22]:

```
# check value counts in hour_of_day column  
df.hour_of_day.value_counts()
```

Out[22]:

```
19    3118  
18    3053  
21    2987  
20    2877  
22    2741  
17    2601  
14    2556  
12    2440  
23    2431  
13    2404  
15    2377  
11    2360  
8     2360  
9     2318  
10    2226  
16    2209  
7     1889  
0     1770  
1     1353  
6     1138  
2      970  
3      737  
4      542  
5      537
```

Name: hour\_of\_day, dtype: int64

In [23]:

```
# check value counts in year column  
df.year.value_counts()
```

Out[23]:

```
2016    49994
```

Name: year, dtype: int64

In [24]:

```
# check value counts in month column  
df.month.value_counts()
```

Out[24]:

```
3     8757  
4     8664  
5     8568  
2     8234  
1     7892  
6     7879
```

Name: month, dtype: int64

In [25]:

```
# check value counts in day_of_month column  
df.day_of_month.value_counts()
```

Out[25]:

```
5      1787  
16     1781  
6      1761  
13     1726  
17     1699  
15     1699  
20     1696  
4       1696  
7       1686  
12      1686  
9       1680  
21      1677  
14      1674  
11      1667  
3       1665  
19      1661  
27      1635  
22      1630  
18      1626  
26      1625  
8       1622  
10      1618  
2       1609  
1       1605  
28      1560  
25      1558  
29      1544  
24      1531  
23      1439  
30      1312  
31       839
```


Name: day\_of\_month, dtype: int64

## Handle Outliers



In [26]:


```
# check outliers by box plot
df.boxplot()
plt.xticks(rotation = 'vertical')
plt.show()
```



**We can see in trip duration column there are much amount of outliers present, so first we have to drop them**

In [27]:

```
# find outliers in trip duration column by box plot
sns.boxplot(df.trip_duration)
plt.show()
```



## Remove Outliers By IQR

In [28]:

```
# remove outliers of trip duration column
q1 = df['trip_duration'].quantile(0.25)
q3 = df['trip_duration'].quantile(0.75)
iqr = q3-q1
first = q1-(1.5*iqr)
third = q1+(1.5*iqr)
```

In [29]:

```
df[df['trip_duration'] < third].shape[0]/df.shape[0]
```

Out[29]:

0.8560027203264392

**If we remove outliers from trip duration column, 15 % of data will be remove**

In [30]:

```
# remove outliers and create new dataframe
df2 = df[(df['trip_duration'] < third) & (df['trip_duration'] > first)]
```

In [31]:

```
df2.shape
```

Out[31]:

```
(42795, 13)
```

In [ ]:

## Question & Answers:-

**1. Choose the most suitable evaluation metric and state why you chose it**

### Answer:-

We Have 6 Evaluation Metrics For Regression Problems, But For Given Dataset I Am Going To Use MSE Evaluation Metric Because Our Dataset Comes With Outliers And Many Lower And Higher Range Values, And MSE Can Handle That Things And Perform Very Well, Thats Why I Am Going To Use MSE, But We Can Use Other Evaluation Metrics As Well

**2. Build a benchmark model for the given dataset.**

### Answer:-

In [32]:

```
# import shuffle to shuffle the data
from sklearn.utils import shuffle

# shuffle the data with random state 42
df3 = shuffle(df2, random_state= 42)

# creating 4 divisions for train and test
div = int(df3.shape[0]/4)

# creating train and test set
train = df3.loc[:3*div+1,:]
test = df3.loc[3*div+1:]
```

## Simple Mean (Mean Of trip\_duration)

### Train Error

In [33]:

```
# create new column in train dataset by mean of trip duration of test dataset
train['simple_mean'] = test['trip_duration'].mean()
```

In [34]:

```
# calculate mean squared error
from sklearn.metrics import mean_squared_error as mse
```

In [35]:

```
train_simple_mean_error = mse(train['trip_duration'],train['simple_mean'])
train_simple_mean_error
```

Out[35]:

111105.56163509692

### Test Error

In [36]:

```
# create new column in test dataset by mean of trip duration of train dataset
test['simple_mean'] = train['trip_duration'].mean()
```

In [37]:

```
test_simple_mean_error = mse(test['trip_duration'],test['simple_mean'])
test_simple_mean_error
```

Out[37]:

110060.71173327156

## Mean Trip Duration With Vendor Id

In [38]:

```
ven_id = pd.pivot_table(train, values='trip_duration', index = ['vendor_id'], aggfunc=np.me
ven_id
```

Out[38]:

trip_duration	
vendor_id	
1	632.166627
2	631.888327

In [39]:

```
test['vend_id_mean'] = 0
for i in train['vendor_id'].unique():
    test['vend_id_mean'][test['vendor_id'] == str(i)] = train['trip_duration'][train['vendo
```

In [40]:

```
ven_id_error = mse(test['trip_duration'], test['vend_id_mean'])
ven_id_error
```

Out[40]:

502544.1163733115

## Mean Trip Duration With Passenger Count

In [41]:

```
pass_count = pd.pivot_table(train, values='trip_duration', index = ['passenger_count'], agg
pass_count
```

Out[41]:

trip_duration	
passenger_count	
1	629.402210
2	644.556359
3	653.268000
4	623.089474
5	616.251836
6	637.826923

In [42]:

```
test['pass_count_mean'] = 0

for i in train['passenger_count'].unique():
    test['pass_count_mean'][test['passenger_count'] == str(i)] = train['trip_duration'][tra
```

In [43]:

```
pass_count_error = mse(test['trip_duration'], test['pass_count_mean'])
pass_count_error
```

Out[43]:

502544.1163733115

## Mean Trip Duration With Day Of Week

In [44]:

```
week_day = pd.pivot_table(train, values='trip_duration', index = ['day_of_week'], aggfunc=n
week_day
```

Out[44]:

trip_duration	
day_of_week	
0	612.766846
1	646.114137
2	652.329050
3	652.897426
4	645.210096
5	625.498370
6	586.921175

In [45]:

```
test['week_day_mean'] = 0

for i in train['day_of_week'].unique():
    test['week_day_mean'][test['day_of_week'] == str(i)] = train['trip_duration'][train['da
```

In [46]:

```
week_day_error = mse(test['trip_duration'], test['week_day_mean'])
week_day_error
```

Out[46]:

502544.1163733115

## Mean Trip Duration With Hour Of Day

In [47]:

```
hour_day = pd.pivot_table(train, values='trip_duration', index = ['hour_of_day'], aggfunc=n
hour_day
```

Out[47]:

trip_duration	
hour_of_day	
0	614.269697
1	610.636029
2	609.330749
3	576.304054
4	648.898477
5	553.540670
6	522.394209
7	584.368946
8	621.260095
9	626.975666
10	654.193267
11	659.826603
12	652.777297
13	638.467301
14	655.341758
15	658.206235
16	640.689133
17	622.133625
18	649.309059
19	623.541340
20	608.386593
21	654.889744
22	650.618526
23	657.047458

In [48]:

```
test['hour_day_mean'] = 0

for i in train['hour_of_day'].unique():
    test['hour_day_mean'][test['hour_of_day'] == str(i)] = train['trip_duration'][train['ho
```

In [49]:

```
hour_day_error = mse(test['trip_duration'], test['hour_day_mean'])
hour_day_error
```

Out[49]:

502544.1163733115

**We Can See Our Simple Mean Model Is Performing Well Comparison Other Benchmark Model**

In [ ]:

### 3. Build a K-Nearest neighbours' model for the given dataset and find the best value of K.

**Answer:-**

In [50]:

```
# map value of store and fwd flag column
df2['store_and_fwd_flag'] = df2['store_and_fwd_flag'].map({'N':0, 'Y':1})
```

In [51]:

df2.head()

Out[51]:

	vendor_id	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	1	1	-73.967026	40.772392	-73.956032	40.76
1	2	1	-73.971046	40.787750	-73.951271	40.77
2	1	2	-73.998093	40.757507	-73.985916	40.74
3	2	1	-73.990631	40.738735	-73.990746	40.73
4	2	1	-73.959160	40.763332	-73.974319	40.75



## Split Dependent And Independent Variable

In [52]:

```
X = df2.drop('trip_duration', axis = 1)
y = df2['trip_duration']
```

In [53]:

```
# split train and test data
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

## Scale Data By Standard Scaller

In [54]:

```
# data scaling by standard scaler
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

In [55]:

```
# import KNN Regressor and MSE Evaluation Metric
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error as mse

knn = KNeighborsRegressor(n_neighbors = 5)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)

score = mse(y_test, y_pred)
score
```

Out[55]:

74396.7812034116

## Elbow Curve For Regressor

In [56]:

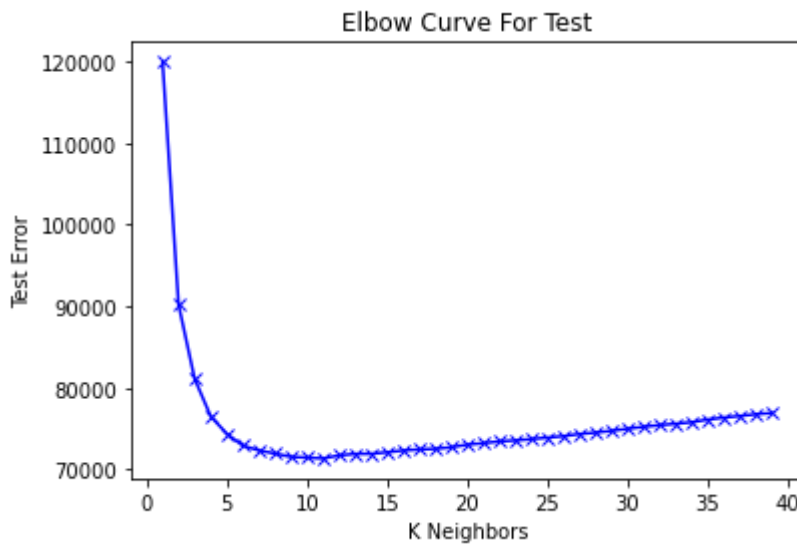
```
def elbow(k):
    test_error = []
    for i in k:
        knn_reg = KNeighborsRegressor(n_neighbors = i)
        knn_reg.fit(X_train,y_train)
        y_pred = knn_reg.predict(X_test)
        score = mse(y_pred,y_test)
        test_error.append(score)
    return test_error
```

In [57]:

```
k = range(1,40)
testt = elbow(k)
```

In [58]:

```
plt.plot(k,testt, 'bx-')
plt.xlabel('K Neighbors')
plt.ylabel('Test Error')
plt.title('Elbow Curve For Test')
plt.show()
```



**Based On Elbow Curve Visualization, We Can See K = 11 Giving Us Very Low Error, So We Will Take 11 As Our K Value**

In [59]:

```
knn_reg = KNeighborsRegressor(n_neighbors = 11)
knn_reg.fit(X_train,y_train)
y_pred = knn_reg.predict(X_test)
knn_train_score = mse(y_train, knn_reg.predict(X_train))
knn_test_score = mse(y_test, y_pred)
print('KNN Train Score:', knn_train_score)
print('KNN Test Score:', knn_test_score)
```

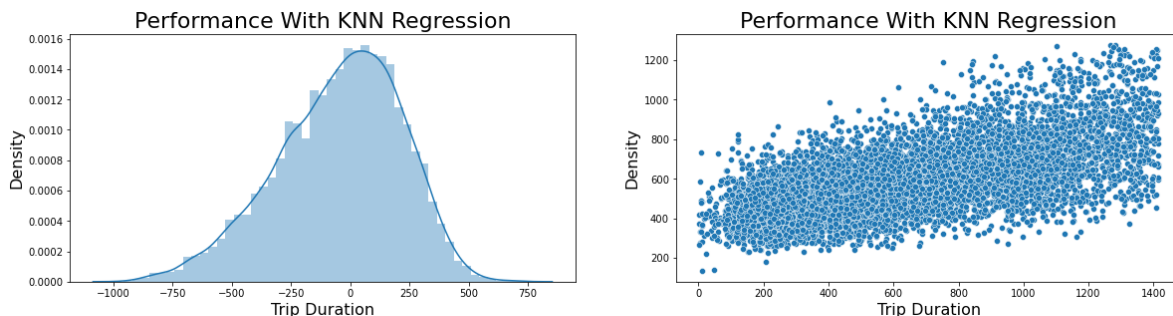
KNN Train Score: 58614.881253989086  
KNN Test Score: 71359.51593846892

In [60]:

```
# visualize the performance of our model
plt.figure(figsize = (20,10))
plt.subplot(2,2,1)
sns.distplot(y_pred-y_test)
plt.title('Performance With KNN Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.subplot(2,2,2)
sns.scatterplot(y_test,y_pred)
plt.title('Performance With KNN Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.show()
```



In [ ]:

## 4. Build a Linear model for the given dataset with regularisation. Attempt to interpret the variable coefficients of the Linear Model.

### Answer:-

In [61]:

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()

lr.fit(X_train,y_train)
y_pred = lr.predict(X_test)
lr_score = mse(y_test,y_pred)
lr_score
```

Out[61]:

109414.47707334299

In [62]:


```
lr.coef_
```

Out[62]:

```
array([-5.19829686e-01,  3.98751720e+00, -1.89271929e+01, -1.16485042e+01,  
       5.65809922e+01, -3.36560477e+01,  1.63220067e+00, -8.32855556e+00,  
       1.56484227e+01,  1.00364161e-13,  7.83150138e+00, -3.80417553e-01])
```

In [63]:

```
xx = X.columns  
yy = lr.coef_  
plt.bar(xx,yy)  
plt.xlabel('Variables')  
plt.xticks(rotation = 'vertical')  
plt.ylabel('Coefficients')  
plt.title('Coefficient Plot')  
plt.show()
```



In [64]:


```
residuals = pd.DataFrame({  
    'fitted_values': y_test,  
    'predicted_values': y_pred})  
  
residuals['residuals'] = residuals['fitted_values'] - residuals['predicted_values']  
residuals.head()
```

Out[64]:

	fitted_values	predicted_values	residuals
35841	692	588.174255	103.825745
22215	872	626.540976	245.459024
6909	709	595.793604	113.206396
8120	319	634.797098	-315.797098
26596	228	581.870220	-353.870220

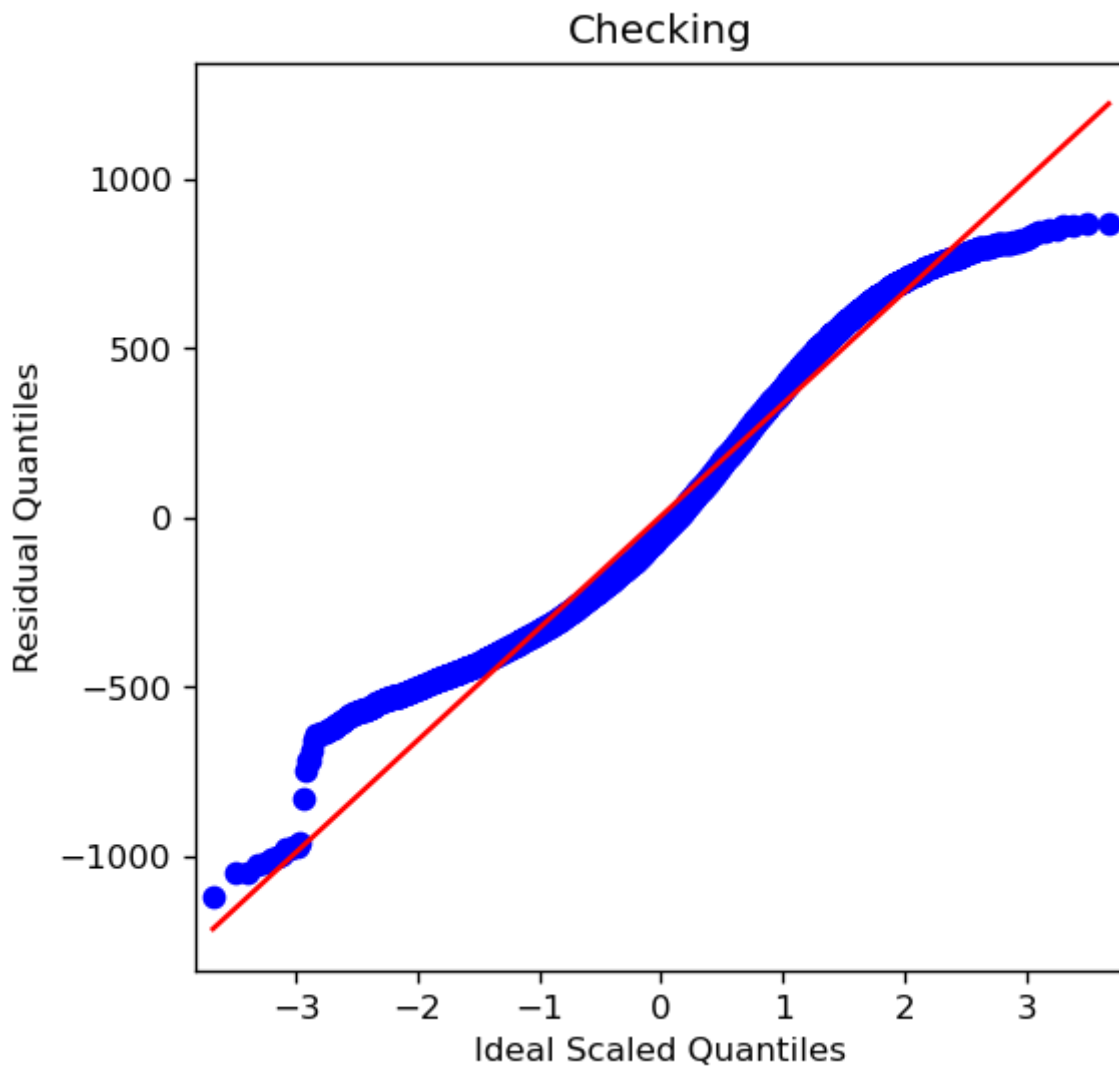
In [65]:

```
plt.hist(residuals.residuals, bins = 100)  
plt.show()
```



In [66]:

```
# importing the QQ-plot from the statsmodels
from statsmodels.graphics.gofplots import qqplot
##Plotting the QQ plot
fig, ax = plt.subplots(figsize=(5,5), dpi = 120)
qqplot(residuals.residuals, line = 's', ax = ax)
plt.ylabel('Residual Quantiles')
plt.xlabel('Ideal Scaled Quantiles')
plt.title('Checking')
plt.show()
```



## We Can See For 60 % Quantiles Its Fit Well

In [67]:

```
variable_coef = pd.DataFrame({  
    'variables': X.columns,  
    'coefficients': lr.coef_  
})  
  
variable_coef.head()
```

Out[67]:

	variables	coefficients
0	vendor_id	-0.519830
1	passenger_count	3.987517
2	pickup_longitude	-18.927193
3	pickup_latitude	-11.648504
4	dropoff_longitude	56.580992

## Chossing variables with sigificance greater than 0.5 (Filtering Significant Features)

In [68]:

```
sig_var = variable_coef[variable_coef.coefficients > 0.5]
```

## Extracting the significant subset do independent Variables

In [69]:

```
subset = df2[sig_var['variables'].values]  
subset.head()
```

Out[69]:

	passenger_count	dropoff_longitude	store_and_fwd_flag	hour_of_day	month
0	1	-73.956032	0	15	1
1	1	-73.951271	0	11	3
2	2	-73.985916	0	19	1
3	1	-73.990746	0	18	2
4	1	-73.974319	0	17	2

In [70]:

```
X_train2,X_test2,y_train2,y_test2 = train_test_split(subset,y,test_size=0.3)
```

In [71]:

```
lr = LinearRegression()
lr.fit(X_train2,y_train2)
y_pred2 = lr.predict(X_test2)
lr_train_score = mse(y_train2, lr.predict(X_train2))
lr_test_score = mse(y_test2, y_pred2)
print('LR Train Score:', lr_train_score)
print('LR Test Score:', lr_test_score)
```

LR Train Score: 109498.6457598188

LR Test Score: 110275.18142869965

## Linear Regression Final Coefficient & Intercept

In [72]:

```
lr.coef_
```

Out[72]:

```
array([ 1.80088622, 684.54472278, 15.26882839,  2.14305419,
        5.6576108  ])
```

In [73]:

```
lr.intercept_
```

Out[73]:

```
51215.61620077867
```


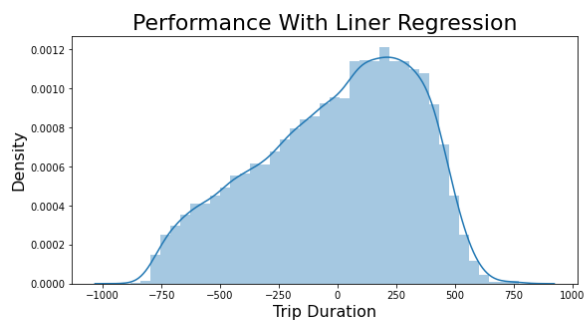


In [74]:

```
# visualize the performance of our model
plt.figure(figsize = (20,10))
plt.subplot(2,2,1)
sns.distplot(y_pred2-y_test2)
plt.title('Performance With Liner Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.subplot(2,2,2)
sns.scatterplot(y_test2,y_pred2)
plt.title('Performance With Liner Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.show()
```



In [ ]:

## Ridge And Lasso Regression

### Ridge Regression

In [75]:

```

from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV
ridge = Ridge()
params_grid = {'alpha' : [0, 1e-8, 1e-5, 1e-3, 1e-2, 1, 2, 4, 5, 10, 20]}
greed_rid = GridSearchCV(ridge,params_grid,cv = 10, n_jobs = 1,scoring = 'neg_mean_squared_

greed_rid.fit(X_train,y_train)
ridge_train_score = mse(y_train, greed_rid.predict(X_train))
y_pred = greed_rid.predict(X_test)
ridge_test_score = mse(y_test,y_pred)

print('Ridge Train Score:', ridge_train_score)
print('Ridge Test Score:', ridge_test_score)
print('Best Alpha Value:', greed_rid.best_params_)

```

Ridge Train Score: 107645.23492198289

Ridge Test Score: 109414.19232039504

Best Alpha Value: {'alpha': 20}

In [76]:



```

# visualize the performance of our model
plt.figure(figsize = (20,10))
plt.subplot(2,2,1)
sns.distplot(y_pred-y_test)
plt.title('Performance With Ridge Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.subplot(2,2,2)
sns.scatterplot(y_test,y_pred)
plt.title('Performance With Ridge Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.show()

```



In [ ]:

## Lasso Regression

In [77]:

```

from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV
lasso = Lasso()
params_grid = {'alpha' : [0, 1e-8, 1e-5, 1e-3, 1e-2, 1, 2, 4, 5, 10, 20]}
greed_lass = GridSearchCV(lasso,params_grid,cv = 10, n_jobs = 1,scoring = 'neg_mean_squared

greed_lass.fit(X_train,y_train)
lasso_train_score = mse(y_train, greed_lass.predict(X_train))
y_pred = greed_lass.predict(X_test)
lasso_test_score = mse(y_test,y_pred)

print('Lasso Train Score:', lasso_train_score)
print('Lasso Test Score:', lasso_test_score)
print('Best Alpha Value:', greed_rid.best_params_)

```

Lasso Train Score: 107645.23039308368

Lasso Test Score: 109414.17059996555

Best Alpha Value: {'alpha': 20}

In [78]:



```

# visualize the performance of our model
plt.figure(figsize = (20,10))
plt.subplot(2,2,1)
sns.distplot(y_pred-y_test)
plt.title('Performance With Lasso Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.subplot(2,2,2)
sns.scatterplot(y_test,y_pred)
plt.title('Performance With Lasso Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.show()

```



In [ ]:

## 5. Build a Decision tree model for the given dataset.

# Attempt to interpret the variable importance.

## Answer:-

In [79]:

```
from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor()

params_grid = {'max_depth': [1,2,3,4,5,6,7,8,9,10,20,30,50,100]}

greed_dtr = GridSearchCV(dtr,params_grid,cv = 10, n_jobs = 1,scoring = 'neg_mean_squared_er
greed_dtr.fit(X_train,y_train)
dtr_train_score = mse(y_train,greed_dtr.predict(X_train))
y_pred = greed_dtr.predict(X_test)
dtr_test_score = mse(y_test,y_pred)

print('DTR Train Score:', dtr_train_score)
print('DTR Test Score:', dtr_test_score)
print('Best Max Depth Value:', greed_dtr.best_estimator_)
```

DTR Train Score: 66537.84105830992

DTR Test Score: 72348.65889447462



Best Max Depth Value: DecisionTreeRegressor(max\_depth=10)

In [80]:

```
# visualize the performance of our model
plt.figure(figsize = (20,10))
plt.subplot(2,2,1)
sns.distplot(y_pred-y_test)
plt.title('Performance With Decision Tree Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.subplot(2,2,2)
sns.scatterplot(y_test,y_pred)
plt.title('Performance With Decision Tree Regression',fontsize = 22)
plt.xlabel('Trip Duration',fontsize = 16)
plt.ylabel('Density',fontsize = 16)

plt.show()
```



## Feature Impotrance Of Decision Tree Regressor

In [81]:

```
feature_importance = pd.DataFrame({  
    'Variables': X.columns,  
    'Feature Importance': greeed_dtr.best_estimator_.feature_importances_*100  
})  
  
feature_importance = feature_importance.sort_values('Feature Importance', ascending=False)
```

In [82]:

```
feature_importance
```

Out[82]:

	Variables	Feature Importance
3	pickup_latitude	43.118110
5	dropoff_latitude	27.865531
2	pickup_longitude	13.214761
4	dropoff_longitude	9.165591
8	hour_of_day	4.326895
7	day_of_week	1.330911
1	passenger_count	0.334497
11	day_of_month	0.326209
10	month	0.187204
0	vendor_id	0.119607
6	store_and_fwd_flag	0.010683
9	year	0.000000

In [ ]:

## Create Dataframe For Store Models And There Train & Test Scores

In [83]:

```
mse_scores = pd.DataFrame({"Models":["Benchmark Model", "Linear regression", "Ridge Regress  
    "Train_Scores":[train_simple_mean_error,lr_train_score,ridge_train_s  
    "Test_Scores":[test_simple_mean_error,lr_test_score,ridge_test_score
```

In [84]:

```
mse_scores
```

Out[84]:

	Models	Train_Scores	Test_Scores
0	Benchmark Model	111105.561635	110060.711733
1	Linear regression	109498.645760	110275.181429
2	Ridge Regressor	107645.234922	109414.192320
3	Lasso Regression	107645.230393	109414.170600
4	KNN Regressor	58614.881254	71359.515938
5	Decision Tree Regressor	66537.841058	72348.658894

In [ ]:

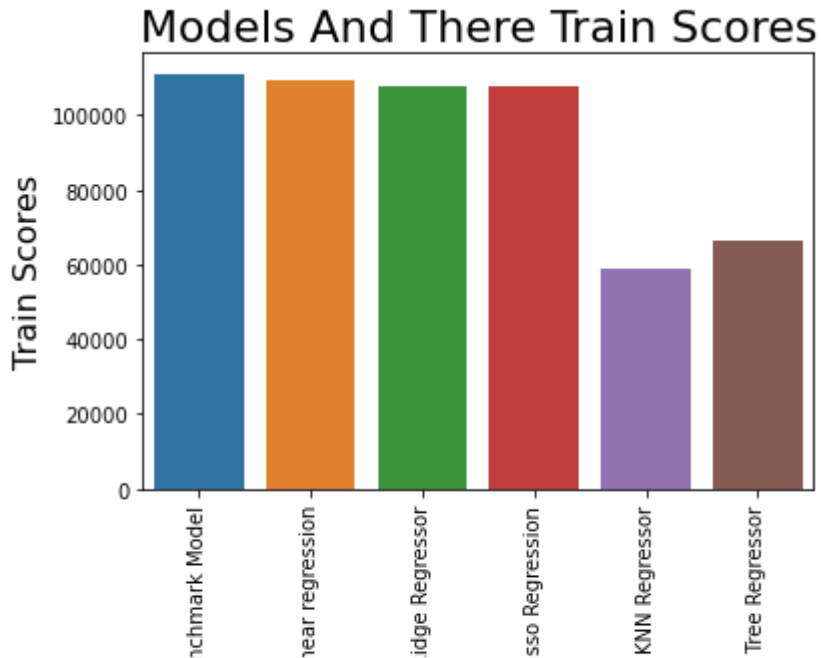
**6. Plot the following Bar plots:**

**0. train score of all the above models.**

**Answer:-**

In [85]:

```
sns.barplot(x = mse_scores.Models,y = mse_scores.Train_Scores)
plt.title('Models And There Train Scores',fontsize = 22)
plt.xlabel('Models',fontsize = 16)
plt.xticks(rotation = 'vertical')
plt.ylabel('Train Scores',fontsize = 16)
plt.show()
```




**1. test (not validation!) score of all the above models.**

**Answer:-**

In [86]:

```
sns.barplot(x = mse_scores.Models,y = mse_scores.Test_Scores)
plt.title('Models And There Test Scores',fontsize = 22)
plt.xlabel('Models',fontsize = 16)
plt.xticks(rotation = 'vertical')
plt.ylabel('Test Scores',fontsize = 16)
plt.show()
```



## 2. Attempt to explain the observations from the plots (optional)



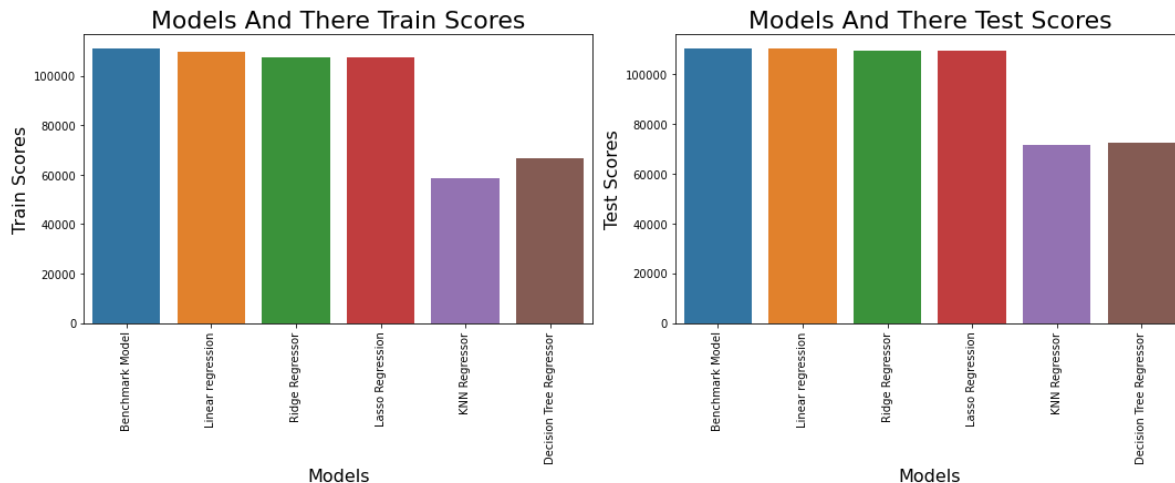
# Answer:-

In [88]:

```
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.barplot(x = mse_scores.Models,y = mse_scores.Train_Scores)
plt.title('Models And There Train Scores',fontsize = 22)
plt.xlabel('Models',fontsize = 16)
plt.xticks(rotation = 'vertical')
plt.ylabel('Train Scores',fontsize = 16)

plt.subplot(2,2,2)
sns.barplot(x = mse_scores.Models,y = mse_scores.Test_Scores)
plt.title('Models And There Test Scores',fontsize = 22)
plt.xlabel('Models',fontsize = 16)
plt.xticks(rotation = 'vertical')
plt.ylabel('Test Scores',fontsize = 16)

plt.tight_layout()
plt.show()
```



1. At the time of training dtr model train very well in compersion with other models and also at the time of testing it's performing very well.
2. After dtr, KNN train well and also at the time of testing it performs well compersion other ML models except dtr model.
3. Benchmark model gave higher mse score at both training and testing time.
4. After evaluating all the models we can say Decision Tree Regressor & KNN Regressor are the best ML model for a given dataset.

In [ ]: